

## *Bayesian Inference*

STAT 211 - 509

2018-10-22

### *Bayes Theorem*

Let  $X$  and  $Y$  be two discrete random variables:

- marginal pmfs  $f_X$  and  $f_Y$  (probabilities of individual values):
  - $P(X = x) = f_X(x)$
  - $P(Y = y) = f_Y(y)$
- conditional pmf of  $Y$ , given the value of  $X$
- $P(Y = y|X = x) = f_{Y|X}(y|x)$

**Bayes Theorem:**

$$P(X = x|Y = y) = \frac{P(Y = y|X = x)P(X = x)}{\sum_x P(Y = y|X = x)P(X = x)}$$

### *Posterior Distributions*

Suppose we want to estimate the unknown parameter of some distribution  $f$ . We obtain a random (IID) sample of observations ( *data* ) from  $f$ , resulting in a likelihood of:


$$L(\theta) = f(\text{data}|\theta)$$

Suppose we are willing to treat the parameter  $\theta$ , a fixed quantity in the population, nevertheless as if it were a random variable, with **prior** distribution  $g$ . Then an alternative strategy for inference could be based on an application of Bayes Theorem:

$$\begin{aligned} \text{posterior} &= h(\theta|\text{data}) \\ &= \frac{f(\text{data}|\theta)g(\theta)}{\sum_{\theta} f(\text{data}|\theta)g(\theta)} \\ &= \frac{\text{likelihood} \times \text{prior}}{\sum \text{likelihood} \times \text{prior}} \end{aligned}$$


### *Illustration*

We have prior beliefs about the vote share  $p$  a candidate will receive in an election:




$p$

We poll one voter in their district to get a sense of voter intention, and update our beliefs:



$p$

We poll many voters and use the combined information to update our beliefs:



### Example: Billiards

- See billiards.R:
- Prior to playing any games, how confident are you that you are the better player in this matchup?
- After winning the first 3 games in a row, now how confident are you?
- After then losing the next 2 games (for a total record of 3 wins and 2 losses), now how confident are you?

Suppose you meet a colleague to discuss business over games of billiards. Suppose you have never played billiards with this particular person before and do not have much confidence in your billiards skills. Let  $p$  be the probability that you win any one game with your colleague.

Going into your first game, suppose your “prior” probabilities for the value of  $p$ , call them

`.pull-left[`  $g(p)$ , are as given in Table 2. `]`

$p$	0.0	0.1	0.2	0.3	0.4	0.5	0.6	0.7	0.8	0.9	1.0
$g(p)$	0.00	0.10	0.25	0.20	0.15	0.10	0.05	0.05	0.05	0.05	0.00

`.pull-right[` `]`

### Frequentist vs. Bayesian Inference

- Frequentist:
  - Probability refers to limiting relative frequencies.
  - Parameters are fixed, unknown constants.

- Statistical procedures designed to have long-run frequency properties.
- Bayesian:
  - Probability described degree of belief, not limiting frequency.
  - We can make probability statements about parameters.
  - We make inferences about a parameter by producing a probability distribution for it.

### *Conjugate Priors*

- When prior and posterior are in same family, the prior is said to be *conjugate* with respect to the model.
- With conjugate priors, can easily draw samples directly from posterior.
- Example:
  - Binomial likelihood and beta prior: posterior also binomial
  - Normal likelihood and normal prior: posterior also normal

### *Functions of Parameters*

- How to make inference about a function  $\tau = g(\theta)$
- Recall how we solved the problem when the density of  $X$  was given as  $f_X$  and we found out density for  $Y = g$ . We will apply the same reasoning here.
- The posterior CDF for  $\tau$  is

$$H(\tau|x^n) = P(g(\theta) \leq \tau) = \int_A f(\theta|x^n) d\theta$$

Where  $A = \{\theta : g(\theta) \leq \tau\}$

- The posterior density is

$$h(\tau|x^n) = H'(\tau|x^n)$$

### Simulation

- The posterior can be approximated by simulation
- If we draw  $\theta_1, \dots, \theta_B \sim p(\theta|x^n)$ , then a histogram of  $\theta_1, \dots, \theta_B$  approximates the density  $p(\theta|x^n)$
- Mean  $\bar{\theta}_n = E(\theta|x^n)$  is

$$\frac{1}{B} \sum_{j=1}^B \theta_j$$

- Let  $\tau_i = g(\theta_i)$ , then  $\tau_1, \dots, \tau_B$  is a sample from  $f(\tau|x^n)$
- This avoids the need for any analytical calculation.

### Credible intervals

- Once we have  $p(\theta|X)$  we can create intervals into which  $\theta$  falls with a certain probability—a **credible interval**.
- Unlike a confidence interval, for a  $1 - \alpha$  credible interval we can say that with probability  $1 - \alpha$   $\theta$  falls in the interval.
- The posterior  $1 - \alpha$  interval can be approximated by  $(\theta_{\alpha/2}, \theta_{1-\alpha/2})$  where  $\theta_{\alpha/2}$  is the  $\alpha/2$  sample quantile of  $\theta_1, \dots, \theta_B$

### Flat Priors

- In case of a more complicated problem where there are many parameters, finding prior  $f(\theta)$  seems impractical.
- An alternative is to define some sort of “noninformative prior”.
- Flat prior  $f(\theta) \propto \text{constant}$  can be used as a noninformative prior.
- Flat priors are not invariant.
- Unfettered use of flat priors raises some questions

### Improper Priors

- If Flat prior  $f(\theta) \propto c$  where  $c > 0$  is a constant, then  $\int f(\theta)d\theta = \infty$
- In usual sense this is not a real probability density. Such priors are called **Improper Prior**
- The Jeffrey’s rule for creating a (invariant) prior:  $f(\theta) \propto I(\theta)^{1/2}$ , where  $I(\theta)$  is the Fisher information
- Improper priors are not a problem as long as the resulting posterior is a well defined probability distribution